

AI Knowledge Management System

A comprehensive, cloud-native knowledge management system built with MCP (Model Context Protocol) servers, featuring AI-powered document processing, semantic search, and knowledge graph capabilities.

Architecture Overview

The system consists of multiple specialized MCP servers orchestrated through a central workflow engine:

- **Phi4 MCP Server:** AI classification and entity extraction using Phi4 model
- **Azure SQL MCP Server:** Document storage and metadata management
- **GraphRAG MCP Server:** Knowledge graph construction and relationship mapping
- **Search MCP Server:** Semantic and vector search capabilities
- **Orchestrator:** Central workflow coordination and API gateway
- **Astro Frontend:** Modern, responsive user interface

Azure Resources

Deployed on Azure with the following services:

Service	Azure Resource Name	Purpose
Static Web App	knowledge-system-ui	Frontend hosting
App Service	phi4-mcp-server	AI processing server
App Service	knowledge-base-sql-server	Database operations
App Service	knowledge-base-graphrag-server	Knowledge graph operations
App Service	knowledge-base-search-server	Search operations
App Service	knowledge-base-orchestrator	Central orchestrator
SQL Database	knowledge-base	Document and metadata storage
SQL Server	knowledge-sql	Database server
AI Foundry	knowledge-ai-foundry	AI model hosting
Storage Account	knowledgestorageacct	File storage
App Service Plan	mcp-app-service-plan	Hosting plan

Features

Core Capabilities

- **AI-Powered Document Processing:** Automatic classification and content analysis
- **Entity Extraction:** Identify people, organizations, concepts, and relationships
- **Knowledge Graph:** Build and visualize connected knowledge networks
- **Semantic Search:** Advanced search using AI embeddings and natural language
- **Multi-format Support:** PDF, DOC, DOCX, TXT, and Markdown files
- **Real-time Processing:** Live status updates and progress tracking

Advanced Features

- **Batch Processing:** Handle multiple documents simultaneously
- **Graph Visualization:** Interactive knowledge network exploration
- **AI Insights:** Generate insights and summaries from knowledge base
- **Search Analytics:** Track and optimize search patterns
- **Workflow Orchestration:** Coordinated multi-step document processing

Repository Structure

knowledge-system/

| — .github/

| — workflows/

| — deploy-frontend.yml

| — deploy-mcp-servers.yml

| — deploy-orchestrator.yml

| — test-integration.yml

| — mcp-servers/

| — phi4-server/

| — package.json

| — server.js

| — web.config

| — sql-server/

| — package.json

| — server.js

| — web.config

| — graphrag-server/

| — package.json

| — server.js

| — web.config

| — search-server/

| — package.json

| — server.js

| — web.config

| — orchestrator/

| — package.json

| — server.js

| — web.config

| — lib/

| — mcpClient.js

| — workflowEngine.js

| — frontend/

| — package.json

| — astro.config.mjs

| — tsconfig.json

| — src/

| — layouts/

| — Layout.astro

| — pages/

| — index.astro

| — upload.astro

| — search.astro

| — public/

```
|   └─ favicon.svg
|   └─ database/
|       └─ schema.sql
|       └─ seed-data.sql
|   └─ docker-compose.yml
|   └─ README.md
|   └─ .gitignore
```

Setup and Deployment

Prerequisites

- Azure subscription with appropriate permissions
- Node.js 18+ installed locally
- Git for version control
- Azure CLI (for local development)

Azure Services Configuration

1. Create Resource Group:

```
bash

az group create --name mcp-knowledge-system --location "East US 2"
```

2. Deploy Azure Resources (use Azure Portal or ARM templates):

- App Service Plan: `mcp-app-service-plan`
- App Services: 5 services for MCP servers and orchestrator
- Azure SQL Database: `knowledge-sql` server with `knowledge-base` database
- Static Web App: `knowledge-system-ui`
- Storage Account: `knowledgestorageacct`
- AI Foundry: `knowledge-ai-foundry`

3. Configure GitHub Secrets:

```
AZURE_STATIC_WEB_APPS_API_TOKEN
AZURE_PHI4_MCP_PUBLISH_PROFILE
AZURE_SQL_MCP_PUBLISH_PROFILE
AZURE_GRAPHRAG_MCP_PUBLISH_PROFILE
AZURE_SEARCH_MCP_PUBLISH_PROFILE
AZURE_ORCHESTRATOR_PUBLISH_PROFILE
```

Database Setup

1. Connect to Azure SQL Database

2. Run schema creation:

```
sql  
  
-- Execute database/schema.sql
```

3. Load sample data (optional):

```
sql  
  
-- Execute database/seed-data.sql
```

Environment Variables

Set these in Azure App Service Configuration:

For SQL Server MCP:

```
AZURE_SQL_SERVER=knowledge-sql.database.windows.net  
AZURE_SQL_DATABASE=knowledge-base  
AZURE_SQL_USERNAME=your_username  
AZURE_SQL_PASSWORD=your_password
```

For Search Server MCP:

```
AZURE_SEARCH_ENDPOINT=https://your-search-service.search.windows.net  
AZURE_SEARCH_API_KEY=your_search_api_key  
AZURE_SEARCH_INDEX=knowledge-base-index
```

For Orchestrator:

```
PHI4_SERVER_URL=https://phi4-mcp-server.azurewebsites.net  
SQL_SERVER_URL=https://knowledge-base-sql-server.azurewebsites.net  
GRAPHRAG_SERVER_URL=https://knowledge-base-graphrag-server.azurewebsites.net  
SEARCH_SERVER_URL=https://knowledge-base-search-server.azurewebsites.net
```

Local Development

1. Clone the repository:

```
bash
```

```
git clone https://github.com/software-tim/knowledge-system.git
cd knowledge-system
```

2. Install dependencies for each service:

```
bash

# Frontend
cd frontend && npm install

# Orchestrator
cd ../orchestrator && npm install

# Each MCP server
cd ../mcp-servers/phi4-server && npm install
cd ../sql-server && npm install
cd ../graphrag-server && npm install
cd ../search-server && npm install
```

3. Start development servers:

```
bash

# Terminal 1: Frontend
cd frontend && npm run dev

# Terminal 2: Orchestrator
cd orchestrator && npm run dev

# Terminal 3-6: MCP Servers
cd mcp-servers/phi4-server && npm run dev
cd mcp-servers/sql-server && npm run dev
cd mcp-servers/graphrag-server && npm run dev
cd mcp-servers/search-server && npm run dev
```

Deployment

Deployment is automated through GitHub Actions when you push to the `main` branch:

1. **Frontend:** Deployed to Azure Static Web Apps
2. **MCP Servers:** Deployed to respective Azure App Services
3. **Orchestrator:** Deployed to Azure App Service

Monitor deployments in the GitHub Actions tab of your repository.

API Documentation

Orchestrator Endpoints

POST `/api/process-document`

Process a single document with AI analysis.

Request:

```
javascript

// Form data with file upload
{
  file: File,
  title: string,
  options: {
    enable_classification: boolean,
    enable_entities: boolean,
    enable_graph: boolean
  }
}
```

Response:

```
javascript

{
  success: true,
  document_id: "12345",
  processing_results: {
    classification: "Technical Documentation",
    entities_extracted: 15,
    graph_relationships: 8
  }
}
```

POST `/api/search`

Search the knowledge base using various methods.

Request:

```
javascript
```

```
{
  query: "machine learning",
  filters: {
    search_type: "semantic",
    content_type: "documents",
    classification: "Technical Documentation"
  },
  limit: 10
}
```

Response:

```
javascript

{
  success: true,
  results: [
    {
      id: "1",
      title: "ML Introduction",
      content_preview: "Machine learning is...",
      score: 0.95,
      classification: "Technical Documentation"
    }
  ],
  total: 25,
  processing_time: 0.34
}
```

POST `/api/generate-insights`

Generate AI insights from the knowledge base.

Request:

```
javascript

{
  prompt: "What are the main themes in my documents?",
  context_documents: ["1", "2", "3"],
  model: "phi4"
}
```


MCP Server Endpoints

Each MCP server exposes:

- `GET /health` - Health check
- Server-specific tool endpoints under `/tools/`

Configuration

Frontend Configuration (`frontend/astro.config.mjs`)

```
javascript

export default defineConfig({
  output: 'static',
  adapter: node({
    mode: 'standalone'
  }),
  server: {
    port: 3000
  }
});
```

Database Connection

Configure in Azure App Service settings or local `.env` file:

```
AZURE_SQL_SERVER=knowledge-sql.database.windows.net
AZURE_SQL_DATABASE=knowledge-base
AZURE_SQL_USERNAME=your_username
AZURE_SQL_PASSWORD=your_password
```

Testing

Run Health Checks

```
bash
```

Test all services

`curl https://knowledge-base-orchestrator.azurewebsites.net/health`

Test individual MCP servers

`curl https://phi4-mcp-server.azurewebsites.net/health`

`curl https://knowledge-base-sql-server.azurewebsites.net/health`

`curl https://knowledge-base-graphrag-server.azurewebsites.net/health`

`curl https://knowledge-base-search-server.azurewebsites.net/health`

Test Document Upload

bash

`curl -X POST https://knowledge-base-orchestrator.azurewebsites.net/api/process-document \`

`-F "file=@test-document.pdf" \`

`-F "title=Test Document" \`

`-F "options={\"enable_classification\":true}"`

Monitoring and Analytics

Available Metrics

- Document processing times
- Search query performance
- Entity extraction accuracy
- System health status
- User activity patterns

Database Views

- `vw_document_summary`: Document overview with entity counts
- `vw_entity_relationships`: Knowledge graph relationships
- `vw_popular_searches`: Most frequent search queries

Stored Procedures

- `sp_GetDocumentWithContext`: Retrieve document with full context
- `sp_SearchDocuments`: Advanced document search with ranking
- `sp_GetKnowledgeGraph`: Extract graph data for visualization

Security

Authentication & Authorization

- Azure AD integration for enterprise use
- API key authentication for service-to-service communication
- Role-based access control for sensitive operations

Data Protection

- Encrypted data storage in Azure SQL Database
- Secure file upload validation
- Content sanitization and validation

Network Security

- HTTPS-only communication
- Azure App Service security features
- Private endpoints for database access

Future Enhancements

Planned Features

- **Multi-tenant Support:** Isolated knowledge bases per organization
- **Advanced Analytics:** ML-powered usage analytics and recommendations
- **Real-time Collaboration:** Live document editing and sharing
- **Mobile App:** Native mobile interface for iOS and Android
- **Integration APIs:** Connect with SharePoint, Teams, and other systems

Scaling Considerations

- **Microservices Architecture:** Further decompose for scale
- **Caching Layer:** Redis for improved performance
- **CDN Integration:** Global content delivery
- **Load Balancing:** Distribute traffic across regions

Contributing

1. Fork the repository

2. Create a feature branch: `git checkout -b feature/new-feature`
3. Commit changes: `git commit -m 'Add new feature'`
4. Push to branch: `git push origin feature/new-feature`
5. Submit a Pull Request

Development Guidelines

- Follow TypeScript/JavaScript best practices
- Add tests for new functionality
- Update documentation for API changes
- Ensure Azure deployment works correctly



License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.



Support

Documentation

- [Azure App Service Documentation](#)
- [Astro Documentation](#)
- [MCP Protocol Specification](#)

Issues and Questions

- Create GitHub Issues for bugs and feature requests
- Use GitHub Discussions for questions and community support
- Contact the development team for enterprise support

Troubleshooting

Common Issues:

1. Deployment Failures

- Check GitHub Actions logs
- Verify Azure publish profiles
- Ensure all secrets are configured

2. Database Connection Issues

- Verify connection strings in App Service configuration

- Check Azure SQL firewall rules
- Test database connectivity

3. **File Upload Failures**

- Check file size limits (50MB default)
- Verify supported file formats
- Monitor App Service logs

Built with ❤️ for intelligent knowledge management